
Disentangling the Black Box: Interpretable Protein-Ligand Docking via Sparse Autoencoders

Anonymous Author(s)

Affiliation

Address

email

Abstract

Protein-ligand docking pipelines based on variational autoencoders (VAEs) achieve strong predictive performance but produce opaque latent representations, while the large number of candidate poses generated per complex imposes substantial computational cost as each undergoes expensive energy minimization and refinement regardless of quality. Physics-informed generative approaches like the Friesner Group’s IFD-ML pipeline ground predictions in thermodynamic principles, making learned features more likely to be physically meaningful and interpretability more tractable and valuable than in end-to-end models. We apply sparse autoencoders (SAEs) to the VAE latent space of the IFD-ML pipeline, training on 435,069 poses across 36 protein systems. At a 4Å threshold, SAE-derived features filter 16.5% of poses prior to refinement while retaining 91.1% of native-like conformations, outperforming the pipeline’s own energy score. Fraction-of-variance-explained (FVE) analysis shows the top SAE retains 62% of latent variance while still supporting a high-precision filter. Sparse features activate exclusively on failed docking attempts, each corresponding to a distinct structural failure mode—an interpretable window into what generative models learn about binding physics. These results show SAEs can serve simultaneously as practical computational filters and interpretability tools for building trustworthy drug-discovery pipelines.

1 Introduction

Protein-ligand binding governs essentially all biological processes and underlies the majority of small-molecule drug discovery [Dhakal et al., 2022, Sim et al., 2025]. Predicting how a ligand binds a flexible protein target, the induced fit docking (IFD) problem [Sherman et al., 2006], remains one of computational biochemistry’s hardest tasks, addressed by physics-based pipelines such as IFD-MD (IFD with molecular dynamics) [Miller et al., 2021]. Most end-to-end machine learning (ML) approaches, such as DeepDTA and its graph-based successors, treat binding as a sequence-to-structure mapping without explicitly modeling atomic interactions [Wu et al., 2025]; and are further limited by scarce training data, poor interpretability, and weak generalization to novel proteins or ligands [Lam and Katritch, 2025, Jiang et al., 2026].

Schrödinger and the Friesner Group’s IFD-ML pipeline combines physics-based simulation with variational autoencoder (VAE) and diffusion generative models to sample conformational space, achieving strong performance in cross-docking benchmarks similar to Mansoor et al. [2024] (Konstantinovskiy et al., in preparation). Because it relies on a VAE, each candidate pose is represented by a fixed-dimensional continuous latent vector—the kind of representation a sparse autoencoder (SAE) is designed to decompose into interpretable features (Section 2.2). This creates two problems. First, the pipeline generates thousands of candidate poses per complex, each undergoing expensive Prime energy minimization [Li et al., 2011] and algorithmic refinement regardless of quality—a bottleneck at scale [Wang et al., 2019]. Second, the VAE’s latent vectors are opaque: it is unclear which directions encode binding geometry, energetic favorability, or structural quality; and understanding

Submitted to in the AI for Drug Discovery workshop (NeurIPS 2026).

39 why a model scores a pose highly matters as much as the score itself [Qadri et al., 2025]. This
40 paper addresses both: a practical filter that removes likely-poor poses before refinement, and an
41 interpretability method that opens up what the latent space has learned.

42 Sparse autoencoders (SAEs) decompose opaque neural representations into interpretable, monose-
43 mantic features [Templeton et al., 2024, Gao et al., 2024]. Applied to protein language models, SAEs
44 recover biologically meaningful features corresponding to secondary structure, binding sites, and
45 evolutionary conservation [Adams et al., 2025, Simon and Zou, 2025]. Physics-informed pipelines
46 like IFD-ML are a compelling target: because the latent space is constrained by thermodynamic priors
47 and atomic interactions, learned features are more likely to correspond to real physical quantities—
48 energy terms, geometric constraints, binding physics—than in end-to-end models [Li et al., 2026,
49 Wu et al., 2024]. SAEs have been applied to structure-generation models such as RFdiffusion
50 (FoldSAE) [Zarzecki et al., 2025], RFdiffusion3/RoseTTAFold3 [Yu and Olson, 2026], and protein
51 folding representations (ESMFold/ESM2-3B) [Parsan et al., 2025], but to our knowledge not yet to a
52 physics-informed, molecular-docking generative model.

53 This work applies TopK SAEs to the VAE latent space of the IFD-ML pipeline to serve both goals at
54 once. Concretely, we make three contributions: (1) a practical pre-refinement filter that flags 16.5% of
55 poses for removal while retaining 91.1% of native-like conformations; (2) evidence that the pipeline’s
56 own energy score cannot substitute for this filter, structurally unable to reach the SAE’s operating
57 points; (3) sparse features that isolate distinct, physically interpretable failure modes, from severe
58 rotational misorientation to milder translational displacement, in what generative models learn about
59 binding physics.

60 2 Methods

61 2.1 Dataset

62 We compiled a dataset of 435,069 molecular docking poses across 36 protein-ligand systems generated
63 by the IFD-ML pipeline, split by case (not pose) to prevent leakage. Each pose consists of a 30-
64 dimensional VAE latent vector, ligand RMSD to the native structure (excluding hydrogen atoms),
65 and a free energy prediction from physics-based scoring. We use two RMSD-based pose-quality
66 definitions: native-like at $\text{RMSD} \leq 2\text{\AA}$ (12.3% native-like, following Shim et al. [2022]), and good at
67 $\text{RMSD} \leq 4\text{\AA}$ (36.6% native-like) for the conservative filtering task. Splits are approximately 70%
68 train, 15% validation, 15% test; normalization uses the training split only, and feature/threshold/combo
69 selection for filtering (Section 2.4) draws on pooled train+val, with test reserved for a single final
70 evaluation.

71 2.2 TopK Sparse Autoencoder

72 We train TopK SAEs with $4\times$ expansion (30D input $\rightarrow$ 120D hidden $\rightarrow$ 30D reconstruction) fol-
73 lowing Adams et al. [2025]; the TopK operation retains the K largest activations and zeros the
74 rest, with reconstruction loss as the training objective. The input is mean-centered via a learned
75 pre-encoder bias, initialized to the training-set mean of the normalized activations and subtracted
76 before encoding, following Simon et al. [2026]’s dead-feature mitigation. We trained SAEs across
77 $K \in \{1, 3, 6, 8, 15, 20, 30\}$, 5 seeds per K , for up to 100 epochs using Adam ($\text{lr} = 10^{-3}$, batch size
78 256), selecting the best run per K by validation loss; we report results for $K \in \{3, 8, 15\}$ in the
79 main text; the full diagnostics sweep across all seven K values is in Appendix A.3 (auPR and FVE
80 trends are consistent, though dead-feature rate is not monotonic in K), selecting the operating K per
81 analysis (Section 3).

82 2.3 Pose Quality Classification

83 To assess whether SAE features provide utility beyond raw latents, we trained logistic regression
84 classifiers on (1) 30D normalized VAE latents (L2 penalty, class-balanced) and (2) 120D sparse SAE
85 activations for $K \in \{3, 8, 15, 20, 30\}$ (L1 penalty, class-balanced), fit on train and evaluated on test.
86 auPR is our primary metric, since it handles class imbalance better than ROC-AUC, which we also
87 report for comparison with prior work.

2.4 Threshold-Based Pose Filtering

The paper’s central practical result is a pre-refinement filter, built using a three-way train/select/test split so reported performance cannot be inflated by threshold or feature selection on the test set. For each of the 120 features per SAE, we fit an F1-maximizing activation cutoff and rank candidates by precision on a pooled train+val selection set, requiring recall > 0.02 to exclude degenerate near-zero-recall picks. We take the top 16 candidates, then search all subset unions—evaluated on the same selection set—for the union maximizing recall subject to retaining $\geq 80\%$ of native-like poses; the winning union is applied exactly once to test, which is never touched during fitting or selection. To confirm that precision derives from the SAE decomposition rather than raw latent geometry, we apply the identical procedure to 500 random unit directions and all 60 signed raw latent axes. Finally, as a physics-based baseline, we apply the same discipline directly to the IFD-ML pipeline’s pre-minimization energy score, accounting for its sentinel value (10,000 kcal/mol)—assigned when the ligand is not in the basin that refinement operates in, or the pose is otherwise too structurally poor to score meaningfully.

3 Results

3.1 Reconstruction, Classification, and Sparsity

Does decomposing the latent space into sparse features cost any of the predictive signal present in the raw, dense latents? Table 1 reports auPR for bad-pose classification, FVE, and dead-feature rate, all on the test set. Raw latents match or exceed every SAE K on auPR at both thresholds—the sparse decomposition provides no consistent predictive lift, so its value lies in filtering (Section 3.2) and interpretability (Section 3.4) rather than classification. FVE grows monotonically with K without a matching auPR gain, and dead-feature rate spikes to 25.8% at $K=15$ despite the mean-centering mitigation from Section 2.2 [Simon et al., 2026], likely because the fixed 12-slot auxiliary revival budget does not scale with K ; stronger, K -scaled interventions may be needed.

Table 1: auPR, FVE, and dead-feature rate, raw latents vs. SAE K (test set).

Rep.	auPR (4Å)	auPR (2Å)	FVE	Dead %
Raw latents	0.770	0.855	—	—
SAE $K=3$	0.709	0.858	33.7%	2.5%
SAE $K=8$	0.719	0.821	62.0%	1.7%
SAE $K=15$	0.750	0.846	90.0%	25.8%

3.2 Practical Filtering

We now evaluate the pre-refinement filter from Section 2.4. Table 2 reports the combo filter for $K \in \{3, 8, 15\}$ at both thresholds, alongside the raw-latent-geometry baselines (Section 2.4). $K=8$ and $K=15$ are the reproducible headline operating points, at 4Å and 2Å respectively ($K=3$ ’s combo cleared the 80% good_kept bar on the selection set at 85.2% but dropped to 77.5% on held-out test—a genuine generalization failure rather than a threshold artifact, and one shared by every $K < 8$ we trained, not just $K=3$; Appendix A.5): $K=8$ flags 16.5% of test poses while retaining 91.1% of native-like conformations (precision 0.803); $K=15$ flags 15.7% while retaining 93.0% (precision 0.935). Recall is modest at both (0.21, 0.17)—a deliberate precision-first design choice, since retaining native-like poses matters more than catching every poor one for a pre-refinement filter.

Neither raw-latent-geometry baseline family reaches the good_kept $\geq 80\%$ target at either threshold (0 of 65,535 candidate-union combinations clear it), so Table 2 reports each baseline’s closest achievable point instead—well below what $K=8/K=15$ achieve while flagging fewer poses.

3.3 Comparison to the IFD-ML Pipeline’s Energy Score

Applying the same selection discipline (Section 2.4) to the pipeline’s pre-minimization energy score, 38.4% of test poses at both thresholds are tied at the sentinel value (energy = 10,000 kcal/mol)—occurring when the ligand starts outside the basin that refinement operates in, or the pose is otherwise unsalvageable. Because this is a tied point mass, no energy threshold can flag fewer than 38.4% of poses regardless of tuning—both SAE headline operating points fall below this floor and so are

Table 2: Filtering performance: SAE $K \in \{3, 8, 15\}$ vs. raw-latent-geometry baselines, both thresholds, held-out test set.

Threshold	Source	% Flagged	Precision	Recall	Good Kept
4Å	SAE $K=3$	27.3	0.698	0.301	77.5%
4Å	SAE $K=8$ (headline)	16.5	0.803	0.210	91.1%
4Å	SAE $K=15$	12.3	0.879	0.171	95.9%
4Å	Random directions (best)	33.7	0.605	0.321	63.6%
4Å	Raw latent axes (best)	27.8	0.727	0.319	79.2%
2Å	SAE $K=3$	38.1	0.791	0.353	45.8%
2Å	SAE $K=8$	19.7	0.884	0.204	84.4%
2Å	SAE $K=15$ (headline)	15.7	0.935	0.172	93.0%
2Å	Random directions (best)	33.7	0.824	0.325	59.7%
2Å	Raw latent axes (best)	27.8	0.857	0.280	72.8%

structurally unreachable, not merely outperformed. At its most conservative achievable point (the sentinel-only filter, 38.4% flagged), energy reaches 78.9% good_kept at 4Å and 78.2% at 2Å—both well below the SAE filters despite flagging more than double the poses. An unconstrained best-F1 energy threshold performs even worse, degenerating to flagging 100% of poses.

3.4 Interpretable Features

The SAE’s sparse features are not just useful for filtering—they are individually interpretable, each corresponding to a distinct, physically identifiable failure mode.




Figure 1: Feature 51 (top, thrombin) and Feature 84 (bottom, uPA), both members of the deployed $K=8$ 4Å combo filter (case details and statistics in text). Left/middle: top-activating test poses, docked pose (orange) vs. reference (green); the colored backbone traces in Feature 51’s poses reflect the pipeline’s induced-fit receptor-loop sampling, present in both good and bad poses for this case (see below). Right: activation distribution (active poses only, log-scale count) split by the 2Å native-pose criterion.

Table 2’s $K=8$ combo filter is a union of 16 candidate features (Section 2.4); we highlight two with sharply contrasting structural signatures for detailed interpretation. Feature 51 fires on 711 test poses, 99.3% of which come from a single thrombin case pair [Di Cera et al., 1995] (PDB 1AE8→1D9I [De Simone et al., 1997, Krishnan et al., 2000], 706/711; a second thrombin pair accounts for the remaining 5), at precision 0.809 with a sentinel rate of 86.4%. This is not indiscriminate case detection: the feature fires on just 9.5% of that case’s poses and 10.8% of its bad poses. Under the stricter, standard 2Å native-pose criterion, precision is perfect: all 711 flagged poses are non-native (0 false positives). Its highest-activating poses (Figure 1) show a case-specific rotational failure:

the ligand’s sterically bulky substituent, which should point toward Leu99, instead points toward Ala190—both residues lining the S1/S2 subsites adjacent to the catalytic triad [Bode et al., 1989]. This misorientation is severe enough that most flagged poses fall outside any refinable basin.

Feature 84 tells a different story. It fires on 827 poses, 94.3% concentrated in a single urokinase (uPA) case pair [Spraggon et al., 1995] (PDB 2VIO→1C5W [Frederickson et al., 2008, Katz et al., 2000], 780/827)—the same catalytic serine-protease fold family as thrombin—at precision 0.911 with a sentinel rate of 28.5%, *below* the population baseline; at the 2Å native-pose criterion, precision rises to 0.975 (806/827). Its highest-activating poses show the ligand correctly oriented but translationally displaced within the binding pocket; energies converge to real, physically stable values (clustered around −9,300 to −9,700 kcal/mol) rather than the sentinel value. On the 20 most confidently native poses for this case (RMSD <0.4Å), feature 84’s activation is exactly 0.0000—the feature is silent on good poses.

Together the two features illustrate a rotational-vs-translational split in the geometric errors the SAE decomposes into distinct, monosemantic features: severe rotational misorientation that falls outside any refinable basin (Feature 51), and milder translational displacement that minimization tolerates (Feature 84)—consistent with the monosemanticity hypothesis that sparse features encode distinct failure phenomena [Templeton et al., 2024, Adams et al., 2025, Clark et al., 2026, Bricken et al., 2023].

Case-specific concentration is not unique to Features 51 and 84: across all 16 candidates in the $K=8$ combo (Table 3, test set), 5 concentrate $\geq 90\%$ of firings in a single case pair and 12 concentrate $\geq 50\%$, spanning four structurally and functionally distinct targets—the two proteases already discussed, a phosphatase (PTP1B), and an ionotropic glutamate receptor (GluR5)—rather than one fold family. A mirror-image search (features preferring native-like poses) found nothing comparably precise or concentrated.

Table 3: Case-concentration of the $K=8$ combo filter’s 16 candidate features, by dominant protein target (test set).

Target	# features	Mean dom. case %	Max dom. case %
GluR5 (ion channel)	3	81.4%	100.0%
Thrombin (protease)	8	62.2%	99.3%
PTP1B (phosphatase)	1	98.3%	98.3%
uPA (protease)	4	74.2%	95.5%

4 Conclusion

We have shown that sparse autoencoders can decompose the latent space of a physics-informed docking pipeline into interpretable, case-specific features—high-precision detectors for distinct, rare failure modes. The deployed combo filters remove a substantial share of poor poses while retaining the large majority of native-like conformations at both thresholds (Section 3.2), beating both a raw-latent-geometry baseline and the pipeline’s own energy score, structurally unable to reach these operating points (Section 3.3). SAEs offer a path toward docking pipelines that are not only predictive but mechanistically understood—systems whose failures can be diagnosed, explained, and, eventually, corrected.

That case-specific concentration extends across the $K=8$ combo’s 16 features (Section 3.4) suggests the SAE learns protein-specific failure signatures a global classifier would dilute—these filters are most valuable when the target family is known, complementary to case-agnostic scoring. Recall is modest at both headline operating points, a deliberate precision-first choice rather than a detectability ceiling, and generalization beyond the 36 systems studied here requires further validation. Understanding a generative pipeline’s failures is one step closer to making physics-informed drug discovery interpretable.

Broader Impact This work reduces computational cost and improves interpretability by filtering low-quality poses before refinement and surfacing failure modes the model has not learned. Filtering existing poses, rather than generating new molecules, does not lower the barrier to designing toxic compounds—the field’s primary dual-use concern. A more realistic risk is uneven performance: our filters are case-specific (Section 3.4), so precision may not transfer beyond these 36 systems, and over-reliance could deprioritize rarer disease targets.

References

- Etowah Adams, Liam Bai, Minji Lee, Yiyang Yu, and Mohammed AlQuraishi. From mechanistic interpretability to mechanistic biology: Training, evaluating, and interpreting sparse autoencoders on protein language models. *bioRxiv*, 2025.
- W. Bode, I. Mayr, U. Baumann, R. Huber, S. R. Stone, and J. Hofsteenge. The refined 1.9 Å crystal structure of human alpha-thrombin: interaction with D-Phe-Pro-Arg chloromethylketone and significance of the Tyr-Pro-Pro-Trp insertion segment. *The EMBO Journal*, 8(11):3467–3475, 1989.
- T. Bricken et al. Towards monosemanticity: Decomposing language models with sparse autoencoders. *Transformer Circuits Thread*, 2023. URL <https://transformer-circuits.pub/2023/monosemantic-features/index.html>.
- J. D. Clark, T. J. Dean, and D. Shukla. Learning physical interactions to compose biological large language models (LLMs). *Communications Chemistry*, 9:81, 2026.
- G. De Simone, G. Balliano, P. Milla, C. Gallina, C. Giordano, C. Tarricone, M. Rizzi, M. Bolognesi, and P. Ascenzi. Kinetics, thermodynamics and x-ray crystallographic study of human alpha-thrombin inhibition by alpha-azalysine derivatives. *Journal of Molecular Biology*, 269(4):558–569, 1997.
- A. Dhakal, C. McKay, J. J. Tanner, and J. Cheng. Artificial intelligence in the prediction of protein–ligand interactions: Recent advances and future directions. *Briefings in Bioinformatics*, 23(1):bbab476, 2022.
- Enrico Di Cera, Enriqueta R. Guinto, Alessandro Vindigni, Quoc D. Dang, Youhna M. Ayala, Meng Wuyi, and Alexander Tulinsky. The Na⁺ binding site of thrombin. *Journal of Biological Chemistry*, 270(38):22089–22092, 1995.
- M. Frederickson, O. Callaghan, G. Chessari, M. Congreve, S. R. Cowan, J. E. Matthews, R. McMenamin, D.-M. Smith, M. Vinković, and N. G. Wallis. Fragment-based discovery of mexiletine derivatives as orally bioavailable inhibitors of urokinase-type plasminogen activator. *Journal of Medicinal Chemistry*, 51(2):183–186, 2008.
- L. Gao, T. D. la Tour, H. Tillman, G. Goh, et al. Scaling and evaluating sparse autoencoders. *arXiv preprint arXiv:2406.04093*, 2024.
- J. Jiang, D. Li, G. Wang, and G.-W. Wei. Recent advances in machine learning predictions of protein–ligand binding affinities. *Current Opinion in Structural Biology*, 96:103193, 2026.
- B. A. Katz, R. Mackman, C. Luong, K. Radika, A. Martelli, P. A. Sprengeler, J. Wang, H. Chan, and L. Wong. Structural basis for selectivity of a small molecule, S1-binding, submicromolar inhibitor of urokinase-type plasminogen activator. *Chemistry & Biology*, 7(4):299–312, 2000.
- R. Krishnan, I. Mochalkin, R. Arni, and A. Tulinsky. Structure of thrombin complexed with selective non-electrophilic inhibitors having cyclohexyl moieties at P1. *Acta Crystallographica Section D: Biological Crystallography*, 56:294–303, 2000.
- J. H. Lam and V. Katritch. Navigating structure-based drug discovery with emerging innovations in physics- and knowledge-based approaches. *Npj Drug Discovery*, 2(1):29, 2025.
- Jianing Li, Robert Abel, Kai Zhu, Yixiang Cao, Suwen Zhao, and Richard A. Friesner. The VSGB 2.0 model: A next generation energy model for high resolution protein structure modeling. *Proteins: Structure, Function, and Bioinformatics*, 79(10):2794–2812, 2011.
- Y. Li, R. Zhan, J. Rao, M. Liu, et al. Structure-informed machine learning for drug discovery: A task-centric perspective. *Briefings in Bioinformatics*, 27(1):bbag081, 2026.
- S. Mansoor, M. Baek, H. Park, G. R. Lee, and D. Baker. Protein ensemble generation through variational autoencoder latent space sampling. *Journal of Chemical Theory and Computation*, 20(7):2689–2695, 2024.

239 E. B. Miller, R. B. Murphy, D. Sindhikara, K. W. Borrelli, et al. Reliable and accurate solution to
 240 the induced fit docking problem for protein–ligand binding. *Journal of Chemical Theory and*
 241 *Computation*, 17(4):2630–2639, 2021.

242 Nithin Parsan, David J. Yang, and John J. Yang. Towards interpretable protein structure prediction
 243 with sparse autoencoders. *arXiv preprint arXiv:2503.08764*, 2025.

244 Y. A. Qadri, S. Shaikh, K. Ahmad, I. Choi, et al. Explainable artificial intelligence: A perspective on
 245 drug discovery. *Pharmaceutics*, 17(9):1119, 2025.

246 Woody Sherman, Tyler Day, Matthew P. Jacobson, Richard A. Friesner, and Ramy Farid. Novel
 247 procedure for modeling ligand/receptor induced fit effects. *Journal of Medicinal Chemistry*, 49(2):
 248 534–553, 2006.

249 H. Shim, H. Kim, J. E. Allen, and H. Wulff. Pose classification using three-dimensional atomic
 250 structure-based neural networks applied to ion channel–ligand docking. *Journal of Chemical*
 251 *Information and Modeling*, 62(10):2301–2315, 2022.

252 J. Sim, D. Kim, B. Kim, J. Choi, and J. Lee. Recent advances in AI-driven protein-ligand interaction
 253 predictions. *Current Opinion in Structural Biology*, 92:103020, 2025.

254 E. Simon and J. Zou. InterPLM: discovering interpretable features in protein language models via
 255 sparse autoencoders. *Nature Methods*, 22:2107–2117, 2025.

256 Elana Simon, Etowah Adams, and James Zou. On the relationship between activation outliers and
 257 feature death in sparse autoencoders. *arXiv preprint arXiv:2605.31518*, 2026.

258 Glen Spraggon, Christopher Phillips, Ursula K. Nowak, Christopher P. Ponting, Derek Saunders,
 259 Christopher M. Dobson, David I. Stuart, and E. Yvonne Jones. The crystal structure of the catalytic
 260 domain of human urokinase-type plasminogen activator. *Structure*, 3(7):681–691, 1995.

261 A. Templeton, T. Conerly, J. Marcus, J. Lindsey, et al. Scaling monosemanticity: Extracting
 262 interpretable features from Claude 3 Sonnet. *Transformer Circuits Thread*, 2024. URL <https://transformer-circuits.pub/2024/scaling-monosemanticity/index.html>.
 263

264 E. Wang, H. Sun, J. Wang, Z. Wang, et al. End-point binding free energy calculation with MM/PBSA
 265 and MM/GBSA: Strategies and applications in drug design. *Chemical Reviews*, 119(16):9478–
 266 9508, 2019.

267 M.-H. Wu, Z. Xie, and D. Zhi. A folding-docking-affinity framework for protein-ligand binding
 268 affinity prediction. *Communications Chemistry*, 8(1):108, 2025.

269 Y. Wu, B. Sicard, and S. A. Gadsden. Physics-informed machine learning: A comprehensive review
 270 on applications in anomaly detection and condition monitoring. *Expert Systems with Applications*,
 271 255:124678, 2024.

272 Michael Yu and Matthew L. Olson. VFUSE: Virulent feature understanding with sparse autoencoders.
 273 *arXiv preprint arXiv:2606.10080*, 2026.

274 W. Zarzecki, P. Szymczak, E. Szczurek, and K. Deja. FoldSAE: Learning to steer protein folding
 275 through sparse representations. *arXiv preprint arXiv:2511.22519*, 2025.

A Appendix

A.1 Full case-concentration breakdown for the $K=8$ combo filter

Table 4 extends Table 3 to all 16 individual candidate features (test set), rather than aggregating by protein target. All 16 are members of the deployed $K=8$ combo (Section 3.2).

Table 4: Per-feature case-concentration for all 16 candidates in the $K=8$ 4Å combo filter (test set), sorted by dominant-case %.

Feature	n firing	Precision	Dominant case	Dom. %	# cases
43	196	0.964	GluR5 (3fv1→3fuz)	100.0	1
51	711	0.809	Thrombin (1ae8→1d9i)	99.3	2
98	482	0.263	PTP1B (1c84→2h4g)	98.3	2
2	841	0.735	uPA (2vio→1c5w)	95.5	3
84	827	0.911	uPA (2vio→1c5w)	94.3	3
91	587	0.876	Thrombin (1gj5→1d3d)	83.5	5
100	958	0.902	Thrombin (1gj5→1d3d)	80.2	3
28	25	0.720	GluR5 (3fv1→3fuz)	76.0	2
102	493	0.813	GluR5 (3fv1→3fuz)	68.2	4
53	179	0.480	uPA (2vio→1c5w)	64.8	3
10	594	0.827	Thrombin (1gj5→1d3d)	56.6	5
85	3709	0.942	Thrombin (1ae8→1d9i)	52.1	5
71	1718	0.574	Thrombin (1ae8→1d9i)	49.5	5
12	446	0.863	uPA (2vio→1c5w)	42.2	5
50	360	0.742	Thrombin (1ae8→1d9i)	41.9	5
118	78	0.500	Thrombin (1ae8→1d9i)	34.6	6

A.2 Mirror-image feature search

Section 3.4 reports that a search for the mirror-image case—features that prefer native-like poses rather than flagging bad ones—found nothing comparably precise or concentrated. The best such feature reached only 55.3% precision at 1.6% recall (4Å label) and 22.2% precision at 11.1% recall (2Å label), far below every feature in Table 4. This asymmetry is consistent with the SAE’s TopK sparsity favoring high-precision detectors for rare, structurally distinctive failure modes over the more diffuse signal of "everything went right."

A.3 Full K sweep diagnostics

Table 5 extends Table 1 to all seven trained K values. Dead-feature rate is not monotonic in K : it is worst at $K=1$ (80.8%, where the single-active-feature constraint likely dominates), falls to near zero for $K \in \{3, 6, 8\}$, rises again for $K \in \{15, 20\}$, and partially recovers at $K=30$. The fixed-auxk explanation offered in Section 3.1 for the $K=3/8$ -vs-15 comparison is therefore only a partial account of this pattern, not a complete one—the full sweep suggests at least two distinct regimes (extreme-sparsity starvation at very low K , and revival-budget saturation at high K) rather than a single monotonic cause.

Table 5: auPR, FVE, and dead-feature rate across the full trained K sweep (test set).

K	auPR (4Å)	auPR (2Å)	FVE	Dead %
1	0.660	0.838	−0.2%	80.8%
3	0.709	0.858	33.7%	2.5%
6	0.737	0.867	49.5%	0.8%
8	0.719	0.821	62.0%	1.7%
15	0.750	0.846	90.0%	25.8%
20	0.749	0.846	98.0%	30.8%
30	0.742	0.810	100.0%	16.7%

A.4 Case-concentration for the $K=15$ 2Å combo

Table 6 reports the same case-concentration analysis as Table 4, but for the $K=15$ 2Å headline combo’s 4 features rather than $K=8$ ’s 16. This combo is less case-locked overall—two of its four features

spread across 5–6 distinct cases rather than concentrating like most of $K=8$ ’s candidates—except for Feature 119, which is both small ($n=76$) and perfectly case-specific. Feature 81 is nominally part of the combo but flags zero test poses (Section 3.1), consistent with it being dead on test.

Table 6: Case-concentration for the $K=15$ 2Å combo filter’s 4 features (test set).

Feature	n firing	Precision	Dominant case	Dom. %	# cases
119	76	1.000	Thrombin (1ae8→1d9i)	100.0	1
101	9235	0.948	GluR5 (3fv1→3fuz)	58.6	6
23	1472	0.849	uPA (2vio→1c5w)	53.6	5
81	0	—	(dead on test)	—	0

A.5 Full K sweep: does the selection-set operating point generalize to test?

Section 3.2 shows $K=3$ ’s combo clears the 80% good_kept bar on the selection set (85.2%) but fails on held-out test (77.5%). Table 7 repeats this selection-vs-test comparison for the full K sweep at 4Å. Every combo clears the bar on the selection set (by construction), but only $K \geq 8$ reliably holds up on test— $K \in \{1, 3, 6\}$ all drop below 80% on test, with $K=6$ borderline (79.9%). This is a systematic pattern, not an isolated failure at $K=3$: the selection procedure generalizes reliably once K is large enough, and $K=8$ is the smallest K at which it consistently does.

Table 7: Selection-set vs. test-set good_kept for the winning combo at each trained K (4Å).

K	Selection good_kept	Test good_kept	Test precision	Test recall
1	82.4%	75.6%	0.575	0.191
3	85.2%	77.5%	0.698	0.301
6	82.8%	79.9%	0.754	0.355
8	83.6%	91.1%	0.803	0.210
15	80.8%	95.9%	0.879	0.171
20	84.9%	86.1%	0.621	0.132
30	82.2%	94.7%	0.817	0.137

A.6 Full protein system list

The dataset spans 36 case pairs across 23 distinct protein targets (some targets contribute multiple case pairs, e.g. thrombin, PTP1B, PIM1). Each case ID follows the pattern {target}_{PDB_A}_{target}_{PDB_B}, identifying a cross-docking pair between the two named structures (the same PDB pair appears in Section 3.4 as “PDB A→B”). Two IDs deviate from this pattern and are reproduced verbatim from the pipeline’s case labels rather than corrected: dpp4_1n1m_HIP740_dpp4_2fjp carries an extra ligand/residue identifier (HIP740), and dhodh_4jgd_dhodh_3u2’s second code is 3 characters rather than the standard 4-character PDB ID.

The full list of case pairs: afab_2nnq_afab_3fr4, ask1_4bhn_ask1_4bie, bace1_5clm_bace1_5ezx, cdk2_1g5s_cdk2_2b55, cdk2_1pxj_cdk2_2btr, chk1_2brb_chk1_2hxl, chk1_2c3k_chk1_2cgu, chk1_2c3k_chk1_2cgu, chk1_2cgu_chk1_2ayp, dhodh_4jgd_dhodh_3u2, dpp4_1n1m_HIP740_dpp4_2fjp, dpp4_1n1m_dpp4_2oph, fxa_1wu1_fxa_1xkb, glur5_3fv1_glur5_3fuz, gsk3b_3sd0_gsk3b_4ach, hivrt_2b5j_hivrt_1c0u, hsd1_2irw_hsd1_3ch6, hsp90_1yet_hsp90_2byi, mdm2_3jzk_mdm2_4zyf, pde5_4md6_pde5_3shy, pdk_3sc1_pdk_3nun, pim1_2xiz_pim1_4lmu, pim1_2xj2_pim1_4enx, pim1_3vbt_pim1_4lmu, pim1_4lmu_pim1_4bzo, ptp1b_1c84_ptp1b_1xbo, ptp1b_1c84_ptp1b_2h4g, ptp1b_1no6_ptp1b_1pyn, rho_2esm_rho_2etk, rock1_4yvc_rock1_5wne, throm_1ae8_throm_1d9i, throm_1gj5_throm_1d3d, throm_1k21_throm_1d3d, trypt_2bm2_tryp_3v7t, trypt_5f03_tryp_2za5, upa_2vio_upa_1c5w. The 6 test-set systems (Section 3.4) are throm_1ae8_throm_1d9i, throm_1gj5_throm_1d3d, glur5_3fv1_glur5_3fuz, upa_2vio_upa_1c5w, ptp1b_1c84_ptp1b_2h4g, and hsd1_2irw_hsd1_3ch6.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and Introduction’s three stated contributions (a practical pose-filtering pipeline, a comparison against the IFD-ML pipeline’s energy score, and interpretable failure-mode features) match the paper’s Results (Sections 3.2–3.4) and Conclusion.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The Conclusion explicitly addresses modest recall (0.21/0.17) as a deliberate precision-first design choice rather than a detectability ceiling, and notes that generalization beyond the 36 protein systems and the two highlighted failure modes requires further validation; Results 3.1 separately reports non-trivial dead-feature rates (up to 25.8% at $K=15$) as a limitation of the SAE training.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper presents empirical interpretability and filtering results; it contains no theorems or formal proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Methods (Section 2) specifies the dataset construction, train/validation/test splitting discipline, SAE architecture and training hyperparameters, and the exact feature-selection and combo-filter procedure; the training and analysis code is released (see Question 5).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: The SAE training and analysis code will be released publicly upon publication. No anonymized link is provided at submission time, in keeping with the double-blind review process. The underlying dataset of docking poses cannot be released, as it is generated by the Friesner Group's IFD-ML pipeline, a proprietary, currently unreleased (beta) Schrödinger pipeline.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 2.1 specifies the case-based 70/15/15 train/validation/test split and normalization discipline; Section 2.2 specifies the TopK SAE architecture, K sweep, seeds, optimizer (Adam, lr 10^{-3}), batch size, and epoch budget; Section 2.4 specifies the feature-selection and combo-filter procedure in full.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: auPR, precision, recall, and good_kept are reported as point estimates from the best-of-5-seeds run per K (selected by validation loss) evaluated on a single case-based test split; seed-to-seed and split-to-split variance is not reported.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: SAE training was run on an internal GPU cluster (16 nodes, 8 GPUs each: nodes gpu001–gpu008 with NVIDIA RTX A4000, 16GB VRAM; gpu009–gpu016 with RTX A4500, 20GB VRAM), with each of the 35 training runs ($K \in \{1, 3, 6, 8, 15, 20, 30\}$, 5 seeds) using a single GPU. Given the small model size (30D→120D→30D), per-run compute is modest; however, exact wall-clock time and total compute were not logged and are not reported.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research involves computational analysis of pre-existing molecular docking simulation data, with no human subjects, no personally identifiable information, and no foreseeable conflict with the Code of Ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: See the Broader Impact paragraph preceding the Acknowledgments: it discusses the positive impact of reducing computational cost and improving interpretability in drug discovery, explains why the dual-use risk of novel toxic compound design is low (the method filters existing poses rather than generating molecules), and identifies uneven cross-protein-family performance as the more realistic risk.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The released SAE checkpoints are small interpretability probes trained on pose-quality latents, not generative or de novo molecule-design models, and pose no meaningful risk of misuse; the underlying docking dataset is not released (see Question 5).

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.

- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [No]

Justification: The training dataset is generated by the Friesner Group's IFD-ML pipeline, Schrödinger's induced-fit docking pipeline, which is proprietary and not yet publicly released (currently in beta); no public license exists to cite, and per Question 5 the dataset itself is not released with this paper. The SAE architecture follows the published methodology of Adams et al. [2025], which is cited.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The new asset introduced is the SAE training and analysis code (Question 5), which is documented with a README describing the project structure, setup, and usage; no new dataset or pretrained model checkpoint trained on proprietary data is released.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: This work does not involve crowdsourcing or human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: This work does not involve human subjects research.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: Claude (Anthropic) was used to assist with statistical analysis pipeline development (including identifying and fixing train/test leakage in the feature-selection methodology), baseline comparison implementation, and results interpretation, in addition to writing and editing. All research questions, hypotheses, and scientific judgment calls originated from the authors; the LLM's role was implementation and analysis assistance under author direction, not independent ideation.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.